

Report Common Assignment MPI01

Matrix dot Matrix

Lecturer: Francesco Moscato - fmoscato@unisa.it

Group:

- | | | |
|---------------------|------------|--|
| • Capitani Giuseppe | 0622701085 | g.capitani@studenti.unisa.it |
| • Falanga Armando | 0622701140 | a.falanga13@studenti.unisa.it |
| • Terrone Luigi | 0622701071 | l.terrone2@studenti.unisa.it |

Index

Index	1
Problem description	3
Solution	3
Version 1	3
Version 4	3
Version 2	3
Version 3	4
Version 5	4
Version 6	4
Experimental setup	5
Hardware	5
CPU	5
RAM	6
Software	7
Performance, Speedup & Efficiency	8
Case n° 1	8
Size-1250	9
V1	9
V4	11
Size-2000	12
V1	12
V4	13
Size-5000	14
V1	14
V4	15
Case n° 2	16
Size-1250	17
V1	17
V4	18
Size-2000	19
V1	19
V4	20
Size-5000	21
V1	21

V4	22
Case n° 3	23
Size-1250	24
V1	24
V4	25
Size-2000	26
V1	26
V4	27
Size-5000	28
V1	28
V4	29
Considerations	30
Other considerations	30
Test case	31
API	32
Public Functions	32
Public Functions Documentation	32
Function init	32
Function matrix_dot_matrix	33
Function print_matrices	33
How to run	34
Extras	35
Size-2000	36
V1	36
V4	37
Why is this happening?	38
Size-2000 with serial on pi4	39
V1	39
V4	40
Final considerations	41

Problem description

Parallelize and Evaluate Performances of "Matrix dot Matrix" Algorithm , by using MPI.

Matrix multiplication to use is the "row - column" multiplication, where elements or result matrix is obtained by scalar multiplication of one row of the first Matrix for one column of the second Matrix.

Use MPI DataTypes to model row and column types.

Evaluate performances when reading files with different MPI I/O approaches.

Solution

We implemented an algorithm that evaluates different MPI I/O approaches, referred as Versions, according to what was asked. In this report, only version 1 and 4 are shown since no relevant changes across other versions were observed.

Version 1

Matrix A is read from file in an ordered manner (i.e. using *MPI_File_read_ordered*).

Matrix B is read from file too, but its transposed has been previously written to the corresponding file. This means that the program only has to read matrix B contiguously.

Version 4

Matrix A is read from file with a file view and a displacement computed using processes ranks and *MPI_File_read*.

Matrix B is read as in [Version 1](#).

Version 2

Matrix A is read from file using *MPI_File_seek* and *MPI_File_read_all* using a displacement computed using processes ranks.

Matrix B is read as in [Version 1](#)

Version 3

Matrix A is read from file with a file view and *MPI_File_read_all* using a displacement computed using processes ranks.

Matrix B is read as in [Version 1](#)

Version 5

Matrix A is read from file with a file view and a displacement computed using processes ranks.

Matrix B is read from a file using an MPI Datatype to read the matrix columns. The matrix B file does not contain Matrix B transposed.

Version 6

Matrix A is read from file in an ordered manner (i.e. using *MPI_File_read_ordered*).

Matrix B is read while computing the matrix product, each process reads a column and computes the partial product.

Experimental setup

Hardware

CPU

```
processors          : 8
vendor_id           : GenuineIntel
cpu family          : 6
model               : 142
model name          : Intel(R) Core(TM) i7-8550U CPU @ 1.80GHz
stepping            : 10
microcode           : 0xe0
cpu MHz             : 800.097
cache size          : 8192 KB
physical id         : 0
siblings            : 8
core id             : 0
cpu cores           : 4
apicid              : 0
initial apicid      : 0
fpu                 : yes
fpu_exception       : yes
cpuid level         : 22
wp                  : yes
flags                : fpu vme de pse tsc msr pae mce cx8 apic sep mtrr
pge mca cmov pat pse36 clflush dts acpi mmx fxsr sse sse2 ss ht tm
pbe syscall nx pdpe1gb rdtscp lm constant_tsc art arch_perfmon
pebs bts rep_good nopl xtopology nonstop_tsc cpuid aperfmperf pni
pclmulqdq dtes64 monitor ds_cpl vmx est tm2 ssse3 sdbg fma cx16
xtpr pdcm pcid sse4_1 sse4_2 x2apic movbe popcnt
tsc_deadline_timer aes xsave avx f16c rdrand lahf_lm abm
3dnowprefetch cpuid_fault epb invpcid_single pti ssbd ibrs ibpb
stibp tpr_shadow vnmi flexpriority ept vpid ept_ad fsgsbase
tsc_adjust bmi1 avx2 smep bmi2 erms invpcid mpx rdseed adx smap
clflushopt intel_pt xsaveopt xsavec xgetbv1 xsaves dtherm ida arat
pln pts hwp hwp_notify hwp_act_window hwp_epp md_clear flush_lld
vmx flags : vnmi preemption_timer invvpid ept_x_only ept_ad
ept_lgb flexpriority tsc_offset vtptr mtf vapic ept vpid
unrestricted_guest ple pml ept_mode_based_exec
```

bugs : cpu_meltdown spectre_v1 spectre_v2
spec_store_bypass lltf mds swapgs itlb_multihit srbds
bogomips : 3999.93
clflush size : 64
cache alignment : 64
address sizes : 39 bits physical, 48 bits virtual

RAM

Memory Device

Array Handle: 0x0008
Error Information Handle: Not Provided
Total Width: 64 bits
Data Width: 64 bits
Size: 8192 MB
Form Factor: SODIMM
Set: None
Locator: ChannelA-DIMM0
Bank Locator: BANK 0
Type: DDR4
Type Detail: Synchronous Unbuffered (Unregistered)
Speed: 2400 MT/s
Manufacturer: SK Hynix
Serial Number: 51DC466C
Asset Tag: 9876543210
Part Number: HMA81GS6AFR8N-UH
Rank: 1
Configured Memory Speed: 2400 MT/s
Minimum Voltage: 1.2 V
Maximum Voltage: 1.2 V
Configured Voltage: 1.2 V

Handle 0x000B, DMI type 17, 40 bytes

Memory Device

Array Handle: 0x0008
Error Information Handle: Not Provided
Total Width: 64 bits
Data Width: 64 bits
Size: 8192 MB
Form Factor: SODIMM
Set: None

Locator: ChannelB-DIMM0
Bank Locator: BANK 2
Type: DDR4
Type Detail: Synchronous Unbuffered (Unregistered)
Speed: 2400 MT/s
Manufacturer: SK Hynix
Serial Number: 51DC45B3
Asset Tag: 9876543210
Part Number: HMA81GS6AFR8N-UH
Rank: 1
Configured Memory Speed: 2400 MT/s
Minimum Voltage: 1.2 V
Maximum Voltage: 1.2 V
Configured Voltage: 1.2 V

Software

- Ubuntu 20.04 LTS
- GCC 9
- Swap : 2147479552 bytes ~ 2GB

Performance, Speedup & Efficiency

Case n° 1

In this case study, the main purpose was to analyze the performance of our program in the following build setup:


- The sequential program is compiled without any gcc optimization
- The parallel programs are compiled with the gcc optimization -Ox where $x = 3$

So here we want to highlight the difference between a simple sequential program compared to a parallel one, furthermore the case study is done on multiple sizes that are 1250, 2000, 5000 and with different number of processes (0, 1, 2, 4, 8, 16). In addition, there are many versions of the read of matrix A and B in the files, but we focused on version 1 and version 4. These are the two versions whose results were the most interesting to analyze. In the end, the matrix B was written column wise in order to improve the performance of the dot product.

Size-1250


V1

Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.0128571 4286	9.5457142 86	0.0257142 8571	9.5842857 14	1	1
Parallel	1	0.0104285 7143	9.4494285 71	0.007	9.4668571 43	1.0124041 77	1.0124041 77
Parallel	2	0.008	5.6458571 43	0.0094285 71429	5.6628571 43	1.6924823 41	0.8462411 705
Parallel	4	0.0087142 85714	3.7062857 14	0.0408571 4286	3.7558571 43	2.5518238 18	0.6379559 545
Parallel	8	0.0112857 1429	2.1147142 86	0.0398571 4286	2.166	4.4248779 84	0.5531097 481
Parallel	16	0.1545714 286	2.2588571 43	0.6087142 857	3.0221428 57	3.1713542 9	0.1982096 431



V4


Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.0128571 4286	9.5457142 86	0.0257142 8571	9.5842857 14	1	1
Parallel	1	0.01	9.5174285 71	0.007	9.5347142 86	1.0051990 47	1.0051990 47
Parallel	2	0.0078571 42857	5.6648571 43	0.0085714 28571	5.6811428 57	1.6870348 02	0.8435174 009
Parallel	4	0.0088571 42857	3.629	0.0334285 7143	3.6712857 14	2.6106074 17	0.6526518 542
Parallel	8	0.011	2.1177142 86	0.0437142 8571	2.1724285 71	4.4117840 47	0.5514730 059
Parallel	16	0.1341428 571	2.2421428 57	0.5104285 714	2.8868571 43	3.3199722 88	0.2074982 68



Size-2000


V1

Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.04	39.784285 71	0.0728571 4286	39.897142 86	1	1
Parallel	1	0.0227142 8571	39.365857 14	0.0177142 8571	39.405857 14	1.0124673 27	1.0124673 27
Parallel	2	0.0195714 2857	23.387142 86	0.0207142 8571	23.427142 86	1.7030306 73	0.8515153 363
Parallel	4	0.0227142 8571	14.888714 29	0.0298571 4286	14.941571 43	2.6702106 3	0.6675526 575
Parallel	8	0.0284285 7143	8.6068571 43	0.065	8.7004285 71	4.5856525 95	0.5732065 744
Parallel	16	0.1984285 714	8.6954285 71	0.5747142 857	9.4688571 43	4.2135119 64	0.2633444 978



V4


Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.04	39.784285 71	0.0728571 4286	39.897142 86	1	1
Parallel	1	0.0257142 8571	39.441857 14	0.0177142 8571	39.485428 57	1.0104269 93	1.0104269 93
Parallel	2	0.0202857 1429	23.445714 29	0.0257142 8571	23.491428 57	1.6983702 26	0.8491851 131
Parallel	4	0.0212857 1429	14.897428 57	0.0245714 2857	14.943142 86	2.6699298 29	0.6674824 573
Parallel	8	0.0275714 2857	8.5995714 29	0.0595714 2857	8.6867142 86	4.5928922 66	0.57411153 32
Parallel	16	0.19	8.7114285 71	0.5831428 571	9.4844285 71	4.2065942 67	0.2629121 417



Size-5000


V1

Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.1471428 571	634.59857 14	0.5414285 714	635.28714 29	1	1
Parallel	1	0.14	626.76385 71	0.106	627.01057 14	1.0132000 51	1.0132000 51
Parallel	2	0.1222857 143	367.83942 86	0.1235714 286	368.08528 57	1.7259237 67	0.8629618 834
Parallel	4	0.1404285 714	234.49985 71	0.2325714 286	234.873	2.7048112 93	0.6762028 233
Parallel	8	0.1905714 286	134.371	0.4875714 286	135.049	4.7041232 65	0.5880154 082
Parallel	16	0.5557142 857	138.18028 57	5.056	143.79185 71	4.4181023 56	0.2761313 973



V4

Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.1471428 571	634.59857 14	0.5414285 714	635.28714 29	1	1
Parallel	1	0.1474285 714	626.42685 71	0.1082857 143	626.68214 29	1.0137310 44	1.0137310 44
Parallel	2	0.1234285 714	367.78542 86	0.1312857 143	368.04014 29	1.7261354 64	0.8630677 321
Parallel	4	0.1407142 857	234.49371 43	0.2921428 571	234.92657 14	2.7041945 02	0.6760486 255
Parallel	8	0.192	134.34685 71	0.5224285 714	135.06157 14	4.7036854 09	0.5879606 761
Parallel	16	0.5415714 286	138.05357 14	5.0027142 86	143.59814 29	4.4240623 88	0.2765038 993



Case n° 2

In this case study, the main purpose was to analyze the performance of our program in the following build setup:


- The sequential program is compiled with gcc optimization -Ox where $x = 3$
- The parallel programs are compiled with the gcc optimization -Ox where $x = 3$

So here we want to highlight the difference between the sequential program compared to a parallel one with the same compiler optimization, furthermore the case study is done on multiple sizes that are 1250, 2000, 5000 and with different number of processes (0, 1, 2, 4, 8, 16). In addition, there are many versions of the read of matrix A and B in the files, but we focused on version 1 and version 4. These are the two versions whose results were the most interesting to analyze. In the end, the matrix B was written column wise in order to improve the performance of the dot product.

Size-1250


V1

Version	Processes	ReadFromFile	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.01	9.508571429	0.02428571429	9.542857143	1	1
Parallel	1	0.01042857143	9.449428571	0.007	9.466857143	1.008028007	1.008028007
Parallel	2	0.008	5.645857143	0.009428571429	5.662857143	1.685166498	0.8425832492
Parallel	4	0.008714285714	3.706285714	0.04085714286	3.755857143	2.540793427	0.6351983569
Parallel	8	0.01128571429	2.114714286	0.03985714286	2.166	4.405751202	0.5507189025
Parallel	16	0.1545714286	2.258857143	0.6087142857	3.022142857	3.157645947	0.1973528717



V4


Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.01	9.5085714 29	0.0242857 1429	9.5428571 43	1	1
Parallel	1	0.01	9.5174285 71	0.007	9.5347142 86	1.0008540 22	1.0008540 22
Parallel	2	0.0078571 42857	5.6648571 43	0.0085714 28571	5.6811428 57	1.6797425 07	0.8398712 533
Parallel	4	0.0088571 42857	3.629	0.0334285 7143	3.6712857 14	2.5993229 31	0.6498307 327
Parallel	8	0.011	2.1177142 86	0.0437142 8571	2.1724285 71	4.3927138 82	0.5490892 352
Parallel	16	0.1341428 571	2.2421428 57	0.5104285 714	2.8868571 43	3.3056215 36	0.2066013 46



Size-2000


V1

Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.0357142 8571	39.25	0.11142857 14	39.397142 86	1	1
Parallel	1	0.0227142 8571	39.365857 14	0.0177142 8571	39.405857 14	0.9997788 581	0.9997788 581
Parallel	2	0.0195714 2857	23.387142 86	0.0207142 8571	23.427142 86	1.6816879 08	0.8408439 539
Parallel	4	0.0227142 8571	14.888714 29	0.0298571 4286	14.941571 43	2.6367469 48	0.6591867 369
Parallel	8	0.0284285 7143	8.6068571 43	0.065	8.7004285 71	4.5281841 62	0.5660230 202
Parallel	16	0.1984285 714	8.6954285 71	0.5747142 857	9.4688571 43	4.1607072 81	0.2600442 051



V4


Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.0357142 8571	39.25	0.11142857 14	39.397142 86	1	1
Parallel	1	0.0257142 8571	39.441857 14	0.0177142 8571	39.485428 57	0.9977640 938	0.9977640 938
Parallel	2	0.0202857 1429	23.445714 29	0.0257142 8571	23.491428 57	1.6770858 67	0.8385429 336
Parallel	4	0.0212857 1429	14.897428 57	0.0245714 2857	14.943142 86	2.6364696 66	0.6591174 165
Parallel	8	0.0275714 2857	8.5995714 29	0.0595714 2857	8.6867142 86	4.5353331 03	0.5669166 379
Parallel	16	0.19	8.7114285 71	0.5831428 571	9.4844285 71	4.1538762 78	0.2596172 674



Size-5000


V1

Version	Processes	ReadFrom File	MatrixDot	WriteTonFile	Elapsed	Speedup	Efficiency
Serial	1	0.1485714 286	625.07142 86	0.5414285 714	625.76142 86	1	1
Parallel	1	0.14	626.76385 71	0.106	627.01057 14	0.9980077 802	0.9980077 802
Parallel	2	0.1222857 143	367.83942 86	0.1235714 286	368.08528 57	1.7000446 71	0.8500223 357
Parallel	4	0.1404285 714	234.49985 71	0.2325714 286	234.873	2.6642544 21	0.6660636 052
Parallel	8	0.1905714 286	134.371	0.4875714 286	135.049	4.6335880 2	0.5791985 026
Parallel	16	0.5557142 857	138.18028 57	5.056	143.79185 71	4.3518558 07	0.2719909 88



V4

Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.1485714 286	625.07142 86	0.5414285 714	625.76142 86	1	1
Parallel	1	0.1474285 714	626.42685 71	0.1082857 143	626.68214 29	0.9985308 114	0.9985308 114
Parallel	2	0.1234285 714	367.78542 86	0.1312857 143	368.04014 29	1.7002531 94	0.8501265 972
Parallel	4	0.1407142 857	234.49371 43	0.2921428 571	234.92657 14	2.6636468 78	0.6659117 195
Parallel	8	0.192	134.34685 71	0.5224285 714	135.06157 14	4.6331567 3	0.5791445 912
Parallel	16	0.5415714 286	138.05357 14	5.0027142 86	143.59814 29	4.3577264 73	0.2723579 045



Case n° 3

In this case study, the main purpose was to analyze the performance of our program in the following build setup:


- The sequential program is compiled without gcc optimization
- The parallel programs are compiled without the gcc optimization

So here we want to highlight the difference between the sequential program compared to a parallel one without compiler optimization, furthermore the case study is done on multiple sizes that are 1250, 2000, 5000 and with different number of processes (0, 1, 2, 4, 8, 16). In addition, there are many versions of the read of matrix A and B in the files, but we focused on version 1 and version 4. These are the two versions whose results were the most interesting to analyze. In the end, the matrix B was written column wise in order to improve the performance of the dot product.

Size-1250


V1

Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.0128571 4286	9.5457142 86	0.0257142 8571	9.5842857 14	1	1
Parallel	1	0.009	12.001285 71	0.0072857 14286	12.018142 86	0.7974847 552	0.7974847 552
Parallel	2	0.0081428 57143	7.7404285 71	0.0102857 1429	7.7591428 57	1.2352248 04	0.6176124 02
Parallel	4	0.0101428 5714	4.8527142 86	0.0592857 1429	4.9222857 14	1.9471209 66	0.4867802 415
Parallel	8	0.0137142 8571	3.2358571 43	0.2384285 714	3.4878571 43	2.7479008 81	0.3434876 101
Parallel	16	0.1497142 857	3.2761428 57	0.6918571 429	4.1175714 29	2.3276549 98	0.1454784 374



V4


Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.0128571 4286	9.5457142 86	0.0257142 8571	9.5842857 14	1	1
Parallel	1	0.0115714 2857	12.011	0.0075714 28571	12.030285 71	0.7966798 081	0.7966798 081
Parallel	2	0.0081428 57143	7.3824285 71	0.0094285 71429	7.4001428 57	1.2951487 42	0.6475743 712
Parallel	4	0.012	4.7885714 29	0.0417142 8571	4.8422857 14	1.9792895 92	0.4948223 979
Parallel	8	0.0118571 4286	3.0312857 14	0.0591428 5714	3.1024285 71	3.0892848 92	0.3861606 115
Parallel	16	0.1454285 714	3.1001428 57	0.5697142 857	3.8158571 43	2.5116993	0.1569812 062



Size-2000


V1

Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.04	39.784285 71	0.0728571 4286	39.897142 86	1	1
Parallel	1	0.024	49.828571 43	0.0188571 4286	49.871285 71	0.8000022 916	0.8000022 916
Parallel	2	0.0205714 2857	30.481857 14	0.0295714 2857	30.532 47	1.3067320 47	0.6533660 235
Parallel	4	0.0262857 1429	19.552857 14	0.0362857 1429	19.615285 71	2.0339822 44	0.5084955 61
Parallel	8	0.0315714 2857	12.399142 86	0.093 0.093	12.523285 71	3.1858366 7	0.3982295 837
Parallel	16	0.1995714 286	12.466285 71	0.6597142 857	13.325571 43	2.9940286 67	0.1871267 917



V4


Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.04	39.784285 71	0.0728571 4286	39.897142 86	1	1
Parallel	1	0.0265714 2857	49.490142 86	0.0188571 4286	49.535428 57	0.8054264 192	0.8054264 192
Parallel	2	0.0215714 2857	30.392714 29	0.0362857 1429	30.451	1.3102079 69	0.6551039 844
Parallel	4	0.0245714 2857	19.553428 57	0.0458571 4286	19.623571 43	2.0331234 3	0.5082808 576
Parallel	8	0.0308571 4286	12.481	0.2	12.711857 14	3.1385770 32	0.3923221 289
Parallel	16	0.1795714 286	12.425428 57	0.5655714 286	13.170142 86	3.0293629 53	0.1893351 846



Size-5000


V1

Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.1471428 571	634.59857 14	0.5414285 714	635.28714 29	1	1
Parallel	1	0.1414285 714	780.998	0.1131428 571	781.25285 71	0.8131645 69	0.8131645 69
Parallel	2	0.1275714 286	476.67828 57	0.2064285 714	477.01242 86	1.3318041 73	0.6659020 864
Parallel	4	0.1562857 143	307.55814 29	0.3342857 143	308.04885 71	2.0622934 58	0.5155733 645
Parallel	8	0.2008571 429	194.37028 57	0.7188571 429	195.28971 43	3.2530496 81	0.4066312 102
Parallel	16	0.5798571 429	198.96771 43	6.1087142 86	205.65628 57	3.0890723 36	0.1930670 21



V4

Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.1471428 571	634.59857 14	0.5414285 714	635.28714 29	1	1
Parallel	1	0.1464285 714	776.44214 29	0.1145714 286	776.70314 29	0.8179278 643	0.8179278 643
Parallel	2	0.127	476.45342 86	0.1917142 857	476.77214 29	1.3324753 81	0.6662376 907
Parallel	4	0.1588571 429	305.39028 57	0.3535714 286	305.90257 14	2.0767630 03	0.5191907 507
Parallel	8	0.1995714 286	192.90071 43	0.6962857 143	193.79642 86	3.2781158 43	0.4097644 804
Parallel	16	0.55	198.01671 43	6.538	205.10457 14	3.0973816 84	0.1935863 553



Considerations

Basically we compared the programs with O0 and O3 modifiers and we noticed that the maximum speedups reachable are of 4.70 and 4.60 in the first and second case, while in the third case (Only O0 modifier, both sequential and parallel programs) the maximum speedup is 3.30. Looking at the absolute elapsed times in the tables we can see that the O3 modifier allows us to boost the performances of the parallel programs more than the performance of the sequential ones.

We also noticed that in every trial the max speedup is reached with 8 processes and it varies between 3.30 and 4.70, that is consistent with our system that has 4 physical cores but 8 threads.

Other considerations

The MPI documentation states that:

“

Implementations of data access routines may buffer data to improve performance. This does not affect reads, as the data is always available in the user's buffer after a read operation completes. For writes, however, the MPI_FILE_SYNC routine provides the only guarantee that data has been transferred to the storage device.

”

(MPI: A Message-Passing Interface Standard, n.d., 506)

also:

“

In addition, in order to prevent subsequent shared offset accesses by the same processes from interfering with this collective access, the call might return only after all the processes within the group have initiated their accesses.

...

Once all processes have issued their requests, locations within the file for all accesses can be computed, and accesses can proceed independently from each other, possibly in parallel.

”

(MPI: A Message-Passing Interface Standard, n.d., 525)

This means that the time measured for the read operation is the time taken by a single process to read its portion of the A matrix (without waiting the others) and the time taken for reading the whole matrix B. While in the case of the write it is only the time taken for “initializing the access” and not writing the relative piece of the resulting C matrix.

Finally, it is worth mentioning that the times resulting from the measures of the parallel programs are the maximum times between all the processes, i.e. if we have that one process that takes longer to compute its portion of the resulting C matrix we use its measured time because only when the slower process completes its task the matrix will be fully calculated.

Test case

In the *test* folder there is the file *test.c* which is the file that contains the test cases. We used the decorator pattern to define the test functions.

We tested the matrix dot matrix product several times by asserting that the expected result and the one returned by our function are the same.

API

Public Functions

Type	Name
void	<code>init</code> (double * a, double * b, int rows_a, int columns_a, int rows_b, int columns_b) This function initializes the two matrices A and B needed in the program.
void	<code>matrix_dot_matrix</code> (double * a, double * b, double * c, int n_rows_A, int n_columns_A, int n_rows_B, int n_columns_B) This function makes the dot product between the matrix A and the matrix B that is expected to be allocated columnwise.
void	<code>print_matrices</code> (double * a, double * b, double * c, int rows_a, int columns_a, int rows_b, int columns_b) Prints the matrices to stdout. C should be the result of AxB.

Public Functions Documentation

Function [`init`](#)

```
void init (  
    double * a,  
    double * b,  
    int rows_a,  
    int columns_a,  
    int rows_b,  
    int columns_b  
)
```

Parameters:

- a pointer to the A matrix used in matrix dot product
- b pointer to the B matrix used in matrix dot product
- rows_a number of rows of A
- columns_a number of columns of A
- rows_b number of rows of B
- columns_b number of columns of B

Function `matrix_dot_matrix`

```
void matrix_dot_matrix (  
    double * a,  
    double * b,  
    double * c,  
    int n_rows_A,  
    int n_columns_A,  
    int n_rows_B,  
    int n_columns_B  
)
```

Parameters:

- `a` pointer to the A matrix used in the dot product.
- `b` pointer to the B matrix used in the dot product.
- `c` pointer to the C matrix used to store the result of dot product.
- `n_rows_A` number of rows of A
- `n_columns_A` number of columns of A
- `n_rows_B` number of rows of B
- `n_columns_B` number of columns of B

Function `print_matrices`

```
void print_matrices (  
    double * a,  
    double * b,  
    double * c,  
    int rows_a,  
    int columns_a,  
    int rows_b,  
    int columns_b  
)
```

Parameters:

- `a` pointer to the A matrix used in the dot product.
- `b` pointer to the B matrix used in the dot product.
- `c` pointer to the C matrix used to store the result of dot product.
- `rows_a` number of rows of A
- `columns_a` number of columns of A
- `rows_b` number of rows of B
- `columns_b` number of columns of B

How to run

1. Create a build directory and launch cmake

```
mkdir build
cd build
cmake ..
```
1. Generate executables with `make`
2. To generate measures (TAKES A LOT OF TIME! Our measures are already included so you should skip this step) run `make generate_measures`
3. To extract mean times and speedup curves from them run `make extract_measures`

Results can be found in the `measures/measure` directory, divided by problem size and the gcc optimization option used.

Extras

We have also tried running our program on a raspberry pi cluster. The front end is a pi3, while the worker nodes are 4 pi4.


We immediately noticed that the serial program was slower than the serial program in the previous configuration due to the difference between the single thread performance of the processors.

The parallel programs run much faster with respect to the serial program on the same configuration, as we are using physically different devices. Each device uses up to 4 processes.

Size-2000


V1

Version	Processes	ReadFrom File	Matrix Dot	WriteTo File	Elapsed	Speedup	Efficiency
Serial	1	0.1614285 714	368.2	1.37	369.7314286	1	1
Parallel	1	1.0428571 43	69.63357143	4.985714 286	75.662	4.88661981 7	4.886619 817
Parallel	2	0.2322857 143	35.09757143	13.95585 714	49.28585714	7.50177535 7	3.750887 679
Parallel	4	0.3117142 857	18.78014286	20.29542 857	39.38757143	9.38700750 4	2.346751 876
Parallel	8	1.6	9.413857143	11.07814 286	22.09171429	16.7362036 2	2.092025 452
Parallel	16	1.3014285 71	4.730857143	15.71614 286	21.74885714	17.0000394 1	1.062502 463



V4

Version	Processes	ReadFromFile	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	0.1614285714	368.2	1.37	369.7314286	1	1
Parallel	1	1.022142857	69.62671429	5.296714286	75.94557143	4.868373779	4.868373779
Parallel	2	0.2314285714	35.12085714	12.057	47.40914286	7.798736832	3.899368416
Parallel	4	0.2418571429	18.74328571	16.39142857	35.37642857	10.45134977	2.612837442
Parallel	8	0.7857142857	9.394428571	12.57228571	22.75242857	16.25019621	2.031274526
Parallel	16	1.258714286	4.724857143	14.44942857	20.433	18.09481861	1.130926163



Why is this happening?

At first we did not notice that the sequential program was run on the cluster frontend which is a raspberry pi3. On the contrary the parallel algorithm was executed on the cluster workers, which are raspberry pi4. So the reason why we have very high speedups is that the sequential algorithm is executed on an older device.


As we found out by looking at this very useful [table](#) the pi3 processor, which is a Cortex-A53 has a more shallow pipe, made of 8 stages, than the pi4 processor which is a Cortex-A72. Also the Cortex-A72 also has out-of-order-execution.

We decided to execute the serial program on a pi4 and then compared the results with the previous “mistake”.

Size-2000 with serial on pi4


V1

Version	Processes	ReadFrom File	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	1.0128571 43	69.888571 43	4.17	75.071428 57	1	1
Parallel	1	1.0428571 43	69.633571 43	4.9857142 86	75.662	0.9921946 099	0.9921946 099
Parallel	2	0.2322857 143	35.097571 43	13.955857 14	49.285857 14	1.5231839 91	0.7615919 954
Parallel	4	0.3117142 857	18.780142 86	20.295428 57	39.387571 43	1.9059674 37	0.4764918 593
Parallel	8	1.6	9.4138571 43	11.078142 86	22.091714 29	3.3981712 6	0.4247714 075
Parallel	16	1.3014285 71	4.7308571 43	15.716142 86	21.748857 14	3.4517413 07	0.2157338 317



V4

Version	Processes	ReadFromFile	MatrixDot	WriteToFile	Elapsed	Speedup	Efficiency
Serial	1	1.012857143	69.888571 43	4.17	75.071428 57	1	1
Parallel	1	1.022142857	69.626714 29	5.2967142 86	75.945571 43	0.9884898 771	0.9884898 771
Parallel	2	0.2314285714	35.120857 14	12.057	47.409142 86	1.5834799 8	0.7917399 899
Parallel	4	0.2418571429	18.743285 71	16.391428 57	35.376428 57	2.1220748 28	0.5305187 07
Parallel	8	0.7857142857	9.3944285 71	12.572285 71	22.752428 57	3.2994907 92	0.4124363 49
Parallel	16	1.258714286	4.7248571 43	14.449428 57	20.433	3.6740287 07	0.2296267 942



Final considerations

What we can confirm is that the serial program alone executes much faster on the pi4, rather than on the pi3 due to the above stated reasons. We can see that on the pi3 it took 370 seconds, while on the pi4 it only took 75 seconds. We have nearly **5 of speedup**.

We can also see that the parallel programs execute similarly in both cases of executions, with the sequential on the pi3 and then on the pi4.

This work is licensed under the Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc-sa/4.0/> or send a letter to Creative Commons, PO Box 1866, Mountain View, CA 94042, USA.